


retry_policy: pooled +0.166 (95% CI +0.076 to +0.256), 39 papers / 90 contrasts, I2 = 98%



Published ablations of 'retry_policy', pooled by DerSimonian-Laird random effects over papers (contrasts within a paper averaged first); the square area is the paper's weight and the whiskers are 95% intervals from an APPROXIMATED variance (papers publish none). UPPER BOUND. These are the authors' own published ablations, and an ablation is run and reported in order to show that the component helps: contrasts that came out null are systematically less likely to have been written down with a number. The pooled effect is therefore an upper bound on the component's contribution, not an estimate of it. It is also conditional on the host systems whose authors chose to ablate that component, on the benchmarks they chose, and on one team's implementation of the component rather than on the component itself. Papers publish no standard errors, so the within-contrast variance is approximated (see scripts/analyse_ablations.py::contrast_variance) and the interval inherits that approximation. Mirror image of scripts/analyse_outcomes.py, whose cross-paper estimates are lower bounds (docs/benchmark_caveats.md sec. 5).