


Funnel plot, published harness ablations, all pooled dimensions



Funnel asymmetry here would be evidence of selective reporting, but the standard errors are APPROXIMATED from the two arm scores (papers publish no SEs), so a larger effect mechanically gets a larger SE and part of any asymmetry is induced by the variance model rather than by selection. Read this plot together with the size-based small-study test and the sign shares in `data/analysis/ablation_bias.csv`, not on its own. UPPER BOUND. These are the authors' own published ablations, and an ablation is run and reported in order to show that the component helps: contrasts that came out null are systematically less likely to have been written down with a number. The pooled effect is therefore an upper bound on the component's contribution, not an estimate of it. It is also conditional on the host systems whose authors chose to ablate that component, on the benchmarks they chose, and on one team's implementation of the component rather than on the component itself. Papers publish no standard errors, so the within-contrast variance is approximated (see `scripts/analyse_ablations.py::contrast_variance`) and the interval inherits that approximation. Mirror image of `scripts/analyse_outcomes.py`, whose cross-paper estimates are lower bounds (`docs/benchmark_caveats.md` sec. 5).