


The authors' own arm ranks first in 42 of 75 blocks (56%) against 34.95 expected once the paper's own ablation arms are removed from the chance baseline: exact p = 0.061, paper-clustered p = 0.206, so the overall advantage does not survive. What survives is the 5+-arm subgroup, 14 of 17 against 6.95 expected, p = 4.4e-04



■ observed: own arm ranks first
 ■ chance over the arms that are not the paper's own ablations
 ▨ uncorrected chance (1 / all arms), superseded

This figure is not a result about harness design: it is the measurement of the practice that makes the own-arm-inclusive estimates untrustworthy. A baseline table is written to show the authors' system winning, and in the largest tables it does - but the overall share no longer beats its chance baseline once that baseline stops counting the paper's own switched-off configurations as rival comparators, and the surviving 5+-arm result is best read as evidence that a large arm table in this literature is largely the paper's own ablation table. OWN-ARM CHANCE BASELINE, corrected: a block's chance probability is 1 / (arms that are not the reporting paper's own switched-off configurations), i.e. 1 / max(n_arms - n_ablation_arms, 1), not 1 / n_arms. An ablation arm is the paper's own system with a part removed, not an independently chosen rival: a full system beating its own ablation is the PREMISE of the ablation analysis, so counting those arms as rivals inflates the own-arm result. The uncorrected 1 / n_arms figures are kept beside the corrected ones, labelled, so the correction is auditable. The exact Poisson-binomial tail also treats the blocks as independent Bernoullis, which they are not - they are pseudo-replicated within paper - so a PAPER-CLUSTERED Monte-Carlo p (one null draw per paper, applied to all that paper's blocks) and a one-block-per-paper exact p are reported beside it. OWN-ARM BIAS, the confound that dominates this design: in a paper's own baseline table the authors' system is the one they tuned, while the comparators are usually run at default settings by people who did not build them, and the table exists to show the authors' system winning. 'The paper's own system wins' is therefore close to tautological. The headline estimates here EXCLUDE the authors' own arm and compare comparator arms with each other; the version including that arm is reported separately and is OPTIMISTICALLY BIASED. Scaffold drift (docs/benchmark_caveats.md sec. 5) is reduced but not removed: the block holds one team, one protocol, one benchmark, one split, one metric and one model fixed, so the cross-paper protocol difference is gone, but the rival was run by the reporting paper at some commit while HARNESS-DB codes it at a pinned commit, which may be a different one. That is measurement error on the design variable; it attenuates towards zero, so every effect is a LOWER BOUND on magnitude. Reporting is also self-selected (sec. 3): which rivals appear in a table is the authors' choice. Associational only; no causal reading is supported.