


The comparable set: every key with two or more coded systems, with every system and score
36 keys, 53 systems, 106 observations (one per key x system) | keys with 2 systems: 23, with 3-4: 8, with 5+: 5



n=11	spread 11.0 pp		swe-search 31.0, autocoderover 30.7, infant-agent 30.0, alibaba-lingmaagent 28.3, masai 27.8, sima 27.7, framepilot 27.3, hyperagent 25.3, moatless-tools 24.3, debug-gym 20.4, swe-agent 20.0
n=8	spread 28.3 pp		isea 51.3, lingxi 42.7, orcaloca 41.0, patchpilot 40.2, moatless-tools 38.3, framepilot 37.0, patchwork 37.0, swe-agent 23.0
n=5	spread 29.8 pp		agentscope 63.4, swe-search 62.2, patchpilot 53.6, framepilot 51.2, swe-agent 33.6
n=5	spread 12.0 pp		lingxi 76.9, confucius-code-agent 74.6, moatless-tools 70.8, swe-agent 66.6, mini-swe-agent 64.9
n=5	spread 16.8 pp		autocoderover 38.4, framepilot 36.2, masai 32.6, swe-agent 23.2, mini-swe-agent 21.6
n=4	spread 10.3 pp		coder 28.3, alibaba-lingmaagent 21.2, autocoderover 19.0, swe-agent 18.0
n=4	spread 12.6 pp		augment-swe-bench-verified-agent 65.4, framepilot 63.8, swe-agent 62.4, mini-swe-agent 52.8
n=3	spread 8.1 pp		aguis 15.9, ponder-press 8.7, cradle 7.8
n=3	spread 2.7 pp		framepilot 50.7, debug-gym 50.2, swe-agent 48.0
n=3	spread 7.0 pp		openjiuwen 82.6, live-swe-agent 79.2, mini-swe-agent 75.6
n=3	spread 8.1 pp		evomas 79.1, live-swe-agent 75.4, mini-swe-agent 71.0
n=3	spread 9.4 pp		prometheus 74.4, live-swe-agent 68.4, mini-swe-agent 65.0
n=3	spread 40.0 pp		scribeagent 53.0, weboperator 48.4, webarena-reference-agent 13.1
n=2	spread 17.8 pp		claude-code 46.7, core-agent 28.9
n=2	spread 24.4 pp		claude-code 62.2, core-agent 37.8
n=2	spread 31.6 pp		ai-agent 45.5, react 13.9
n=2	spread 44.2 pp		harness 89.0, nexus 44.9
n=2	spread 0.0 pp		auto-deep-research 55.1, autoagent-2 55.1
n=2	spread 1.8 pp		vlaa-gui 76.3, hippo-agent 74.5
n=2	spread 2.4 pp		os-symphony 65.8, gta 63.4
n=2	spread 1.1 pp		agents-v2 54.2, gta 53.1
n=2	spread 3.2 pp		agents 20.6, agentstore 17.3
n=2	spread 6.9 pp		confucius-code-agent 52.7, live-swe-agent 45.8
n=2	spread 12.2 pp		swe-search 17.0, debug-gym 4.8
n=2	spread 9.4 pp		aegis 30.3, debug-gym 21.0
n=2	spread 0.8 pp		openhands-versa 31.3, zencoder 30.6
n=2	spread 1.2 pp		refact-ai-agent 35.6, openhands-versa 34.4
n=2	spread 10.1 pp		reasoningbank 38.8, mini-swe-agent 28.7
n=2	spread 3.8 pp		reasoningbank 57.4, mini-swe-agent 53.6
n=2	spread 3.4 pp		live-swe-agent 77.4, mini-swe-agent 74.0
n=2	spread 5.5 pp		live-swe-agent 77.4, mini-swe-agent 71.9
n=2	spread 5.0 pp		live-swe-agent 63.0, mini-swe-agent 58.0
n=2	spread 9.8 pp		nemotron-cortexa 68.2, mini-swe-agent 58.4
n=2	spread 19.6 pp		patchpilot 64.6, mini-swe-agent 45.0
n=2	spread 6.8 pp		autocoderover 18.8, swe-agent 12.0
n=2	spread 10.2 pp		browsergym 23.5, webarena-reference-agent 13.3

This is the entire evidence base for any outcome claim in this review. A key is benchmark|split|base-model; a system contributes one observation per key, the median of its rows in that key. Keys holding a single system are excluded because they support no comparison. Standardisation rule: $z = (x - \text{key mean}) / \text{key sd} (\text{ddof}=1)$ for keys with 4 or more systems, standardised average rank $(\text{rank} - (n+1)/2) / \text{sd}(1..n)$ below that, 0 when a key's scores are identical. Raw ranges are percentage points on each benchmark's standard metric and are not comparable across keys. Scaffold drift (docs/benchmark_caveats.md sec. 5): a leaderboard or paper row does not say which commit of the harness produced its number, while HARNESS-DB codes each system at a pinned commit, so the harness coded may not be the harness that scored. That is measurement error on the design variable, not the outcome; no fixed effect absorbs it and it attenuates estimates towards zero. Every effect shown is a LOWER BOUND on magnitude. Reporting is also self-selected (sec. 3): authors report where their system does well, so a within-key comparison is not a randomised one. Associational only; no causal reading is supported.