


The comparable set: every key with two or more coded systems, with every system and score
36 keys, 53 systems, 106 observations (one per key x system) | keys with 2 systems: 23, with 3-4: 8, with 5+: 5



This is the entire evidence base for any outcome claim in this review. A key is benchmark|split|base-model; a system contributes one observation per key, the median of its rows in that key. Keys holding a single system are excluded because they support no comparison. Standardisation rule: $z = (x - \text{key mean}) / \text{key sd} (\text{ddof}=1)$ for keys with 4 or more systems, standardised average rank $(\text{rank} - (n+1)/2) / \text{sd}(1..n)$ below that, 0 when a key's scores are identical. Raw ranges are percentage points on each benchmark's standard metric and are not comparable across keys. Scaffold drift (docs/benchmark_caveats.md sec. 5): a leaderboard or paper row does not say which commit of the harness produced its number, while HARNESS-DB codes each system at a pinned commit, so the harness coded may not be the harness that scored. That is measurement error on the design variable, not the outcome; no fixed effect absorbs it and it attenuates estimates towards zero. Every effect shown is a LOWER BOUND on magnitude. Reporting is also self-selected (sec. 3): authors report where their system does well, so a within-key comparison is not a randomised one. Associational only; no causal reading is supported.