


Why the outcome analysis is small: the attrition from coded system to comparable system



646 of 1,256 coded systems report any benchmark number at all, and only 53 sit in a key shared with another coded system, which is the only configuration that supports a comparison. The gap of 593 systems - reporting a score but never against a peer at the same benchmark, split and base model - is the ceiling on every outcome claim in this review, and it is a finding about the field's reporting practice, not a data-collection failure. Scaffold drift (docs/benchmark_caveats.md sec. 5): a leaderboard or paper row does not say which commit of the harness produced its number, while HARNESS-DB codes each system at a pinned commit, so the harness coded may not be the harness that scored. That is measurement error on the design variable, not the outcome; no fixed effect absorbs it and it attenuates estimates towards zero. Every effect shown is a LOWER BOUND on magnitude. Reporting is also self-selected (sec. 3): authors report where their system does well, so a within-key comparison is not a randomised one. Associational only; no causal reading is supported.