

Design	What it can identify	Headline estimate, in its own units
Cross-paper association 36 comparable keys (benchmark, split, model) over 53 systems	 lower bound: scaffold drift pulls it towards zero	 multi-agent topology +0.88 within-key sd, 95% CI 0.38 to 1.31, on 18 keys
Within-study meta-analysis 75 of the authors' own ablation contrasts, 42 papers	 upper bound: null ablations go unreported	 self-verification +0.139 relative (20 contrasts, 14 papers); discounted for selective reporting +0.069; prediction interval -0.123 to 0.400 crosses zero
corpus-wide 2,466 contrasts from 494 papers, same design; self-verification 680 from 213	 upper bound, as above	 self-verification +0.157 relative, 95% CI 0.132 to 0.183; prediction interval -0.172 to 0.487 crosses zero
Registered compute-matched ablation arms paired by instance; verification against unguided retries at equal calls	 unbiased by design, once it can run	no estimate: the pilot failed its band. 0 of 9 baseline cells lie inside 25–70% on their 95% interval (1 on the point estimate), so no arm comparison is published

Bars: the rung each design reaches; white = lower bound, hatched = upper bound, dark = unbiased. ● estimate with its 95% CI; ◇ discounted for selective reporting; dotted, prediction interval; small arrow, the side of the estimate on which the true value lies.