

Two independent votes

two models vote on every record
27,747 records
agreeing votes, or a fixed rule on their stated reasons, decide

Third-vote tiebreak

no two-vote decision: a third, decisive vote
13,483 records
two of three agreeing votes decide

Inclusive default

no two-of-three decision: forwarded, never excluded
757 records
a missed harness costs more than one extra reading

Title/abstract outcome

voted include	9,114
undecided	757
forwarded by rule	96
to full text	9,967
excluded	17,780

8,435 assessed; 1,532 sought but never retrieved

Tier-1 reader

reads the retrieved record: a verbatim quote per decision step, and a stated confidence
7,227 read
of the 8,435 assessed, 1,135 were read at tier 2 only

Escalation rules

→ tier 2: every tier-1 exclude
 → tier 2: every low-confidence decision
 → tier 2: a hash-selected audit of the other includes
otherwise the tier-1 include stands

Tier-2 reader, decisive

a stronger model reads the record independently; its vote is the decision of record
4,198 decided
no third vote at this stage

Full-text outcome

included	7,085
excluded	1,350
decided at tier 2	4,198
stands at tier 1	4,164
no text, no reading	73
<i>tier 1 alone never excludes</i>	

not escalated: the tier-1 include stands

Coverage the escalation produced

includes read twice
 excludes read twice
 agreement of the two readings on those records
 known harnesses in the reference set kept

Escalation fires on an exclude or on low confidence, so the records read twice are a non-random sample tilted towards excludes and hard cases, and κ describes that sample, not the corpus.

2,073 of 7,085 (29.3%)

1,002 of 1,350 (74.2%)

$\kappa = 0.473$, 73.7% of 3,075

27 of 27 at every stage, search to coding frame