

Two independent votes

two models vote on every record
27,747 records
agreeing votes, or a fixed rule on their stated reasons, decide

Third-vote tiebreak

no two-vote decision: a third, decisive vote
13,483 records
two of three agreeing votes decide

Inclusive default

no two-of-three decision: forwarded, never excluded
757 records
a missed harness costs more than one extra reading

Title/abstract outcome

voted include	9,114
undecided	757
forwarded by rule	96
to full text	9,967
excluded	17,780

8,435 assessed; 1,532 sought but never retrieved

Tier-1 reader

reads the retrieved record: a verbatim quote per decision step, and a stated confidence
7,227 read
of the 8,435 assessed, 1,135 were read at tier 2 only

Escalation rules

→ tier 2: every tier-1 exclude
 → tier 2: every low-confidence decision
 → tier 2: a hash-selected audit of the other includes
otherwise the tier-1 include stands

Tier-2 reader, decisive

a stronger model reads the record independently; its vote is the decision of record
4,198 decided
no third vote at this stage

Full-text outcome

included	7,085
excluded	1,350
decided at tier 2	4,198
stands at tier 1	4,164
no text, no reading	73
<i>tier 1 alone never excludes</i>	

not escalated: the tier-1 include stands

What the escalation changed

tier-1 excludes (every one escalated) overturned to include
 tier-1 includes (low-confidence or sampled) overturned to exclude
 agreement of the two readings on those records
 known harnesses in the reference set kept

670 of 1,532 (43.7%)

140 of 1,543 (9.1%)

$\kappa = 0.473$, 73.7% of 3,075

27 of 27 at every stage, search to coding frame

Escalation fires on an exclude or on low confidence, so the records read twice are a non-random sample tilted towards excludes and hard cases, and κ describes that sample, not the corpus. The sampled includes hold every low-confidence one, so their overturn rate bounds the false-include rate of the unescalated includes from above.